

Name : Kiran Shinde

PRN : 202201040091

Div : A

Batch : GEN AI 2

Practical Experiment 2 :

1. Build a simple AE model for Dimensionality Reduction and Denoising.

Code :

```
#  Practical 1: Autoencoder with MNIST
```

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
```

```
Load MNIST
transform = transforms.ToTensor()
train_dataset = datasets.MNIST(root='./data', train=True, download=True,
                               transform=transform)
train_loader = DataLoader(train_dataset, batch_size=128, shuffle=True)
```

```
Autoencoder
class Autoencoder(nn.Module):
    def __init__(self):
        super(Autoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Flatten(),
            nn.Linear(28*28, 128),
            nn.ReLU(),
```

```
        nn.Linear(128, 32) # compressed representation
    )
    self.decoder = nn.Sequential(
        nn.Linear(32, 128),
        nn.ReLU(),
        nn.Linear(128, 28*28),
        nn.Sigmoid(),
        nn.Unflatten(1, (1, 28, 28))
    )
def forward(self, x):
    z = self.encoder(x)
    return self.decoder(z)

device = "cuda" if torch.cuda.is_available() else "cpu"
model = Autoencoder().to(device)
criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)
```

```
Train AE
for epoch in range(2):    increase epochs for better results
    for imgs, _ in train_loader:
        noisy_imgs = imgs + 0.3*torch.randn_like(imgs)    add noise
        noisy_imgs, imgs = noisy_imgs.to(device), imgs.to(device)

        outputs = model(noisy_imgs)
        loss = criterion(outputs, imgs)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        print(f"Epoch [{epoch+1}/2], Loss: {loss.item():.4f}")

#Visualize original, noisy, denoised
imgs, _ = next(iter(train_loader))
```


```

noisy_imgs = imgs + 0.3*torch.randn_like(imgs)
with torch.no_grad():
    outputs = model(noisy_imgs.to(device)).cpu()

fig, axes = plt.subplots(3, 8, figsize=(12,5))
for i in range(8):
    axes[0,i].imshow(imgs[i][0], cmap="gray"); axes[0,i].axis("off")
    axes[1,i].imshow(noisy_imgs[i][0], cmap="gray"); axes[1,i].axis("off")
    axes[2,i].imshow(outputs[i][0], cmap="gray"); axes[2,i].axis("off")
plt.suptitle("Original | Noisy | Denoised")
plt.show()

```

Output :



2. Generate realistic faces or interpolate between facial features for creative applications in entertainment and design using VAE.

Practical 2: VAE with MNIST

```
import torch.nn.functional as F
```

```
VAE Model
```

```
class VAE(nn.Module):
```

```
    def __init__(self, latent_dim=16):
```

```
        super(VAE, self).__init__()
```

```
        self.fc1 = nn.Linear(28*28, 128)
```

```
        self.fc21 = nn.Linear(128, latent_dim)    mean
```

```
        self.fc22 = nn.Linear(128, latent_dim)    logvar
```

```
        self.fc3 = nn.Linear(latent_dim, 128)
```

```
        self.fc4 = nn.Linear(128, 28*28)
```

```
    def encode(self, x):
```

```
        h1 = F.relu(self.fc1(x))
```

```
        return self.fc21(h1), self.fc22(h1)
```

```
    def reparameterize(self, mu, logvar):
```

```
        std = torch.exp(0.5*logvar)
```

```
        eps = torch.randn_like(std)
```

```
        return mu + eps*std
```

```
    def decode(self, z):
```

```
        h3 = F.relu(self.fc3(z))
```

```
        return torch.sigmoid(self.fc4(h3))
```

```
    def forward(self, x):
```

```
        mu, logvar = self.encode(x.view(-1, 28*28))
```

```
        z = self.reparameterize(mu, logvar)
```

```
        return self.decode(z), mu, logvar
```

```
def loss_function(recon_x, x, mu, logvar):
    BCE = F.binary_cross_entropy(recon_x, x.view(-1, 28*28), reduction='sum')
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD
```

```
vae = VAE(latent_dim=16).to(device)
optimizer = optim.Adam(vae.parameters(), lr=1e-3)
```

```
Train VAE
for epoch in range(2):
    for imgs, _ in train_loader:
        imgs = imgs.to(device)
        recon, mu, logvar = vae(imgs)
        loss = loss_function(recon, imgs, mu, logvar)
```

```
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    print(f"Epoch [{epoch+1}/6], Loss: {loss.item():.4f}")
```

```
#Generate new digits
with torch.no_grad():
    z = torch.randn(8, 16).to(device)
    samples = vae.decode(z).cpu().view(-1, 1, 28, 28)
```


```
fig, axes = plt.subplots(1, 8, figsize=(15,3))
for i in range(8):
    axes[i].imshow(samples[i][0], cmap="gray")
    axes[i].axis("off")
plt.suptitle("Generated Digits")
plt.show()
```

```
Interpolation
with torch.no_grad():
```

```

z1, z2 = torch.randn(1, 16).to(device), torch.randn(1, 16).to(device)
interpolations = []
for alpha in torch.linspace(0, 1, 8):
    z = (1-alpha)*z1 + alpha*z2
    interpolations.append(vae.decode(z).cpu().view(28,28))
fig, axes = plt.subplots(1, 8, figsize=(15,3))
for i in range(8):
    axes[i].imshow(interpolations[i], cmap="gray")
    axes[i].axis("off")
plt.suptitle("Latent Space Interpolation")
plt.show()

```



Code Link :

https://colab.research.google.com/drive/1akrpMdsJnsegwUWrkC2exwu6pgZUuD_j?usp=sharing